

## (Cross-) Validation

### *Simple validation*

```
set.seed(23) # to make the example reproducible

?sample

library(ggplot2)

?diamonds

mydiamonds = diamonds[1:500,]

attach(mydiamonds)

trainingdiamonds = sample(x = 500, size = 200) # 200 observations from 500 in
the df


mylm = lm(data = mydiamonds, subset = trainingdiamonds, x ~ y + z) # fitting a
simple lm

mean ((x - predict(mylm, mydiamonds))[-trainingdiamonds]^2) # mean standard
error

library(boot) # for the cv functions
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



## *LOOCV*

```
myglm = glm(data = mydiamonds, x ~ y + z)
```

```
myglm.error = cv.glm(data = mydiamonds, myglm)
```

```
myglm.error$delta
```

# delta is the cv estimate or the error rate - raw and adjusted

## *Cross Validation K=5*


```
myglm.error2 = cv.glm(data = mydiamonds, myglm, K = 5)
```

# using K to adjust the group number

```
myglm.error2$delta
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



## *Exercise Cross Validation*


Create a model with the faithful dataset

?faithful

- explain waiting with eruptions
- get a visual impression to get an idea of the relationship
- get the MSE using simple cross validation and 5-fold CV
- in the solution I will use a 50/50 split of the dataset for simple CV
- compare the results - which one has a lower error rate?
- what could be possible problems with the CV approach I outlined

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



## 1. simple CV

# simple xy plot

```
plot(faithful$waiting, faithful$eruptions)
```

# linear model explaining the waiting time - training data 1-136

```
mymodel = glm(data = faithful[1:136,], waiting ~ eruptions)
```

# MSE on the second half of the data (validation set)

```
mean((faithful$waiting - predict(mymodel, faithful))[137:272]^2)
```

# Or an alternative way to code it

```
mean((faithful$waiting[137:272] - predict(mymodel, faithful[137:272,]))^2)
```

## 2. 5-fold CV

```
library(boot) # for the cv.glm function
```

# we are going to get a model with the full dataset

```
mymodel2 = glm(data=faithful, waiting ~ eruptions)
```

# cv.glm for 5 fold CV

```
cv.result = cv.glm(data = faithful, mymodel2, K = 5)
```


# the error rate is lower than with standard CV

```
cv.result$delta
```

# if you perform simple CV with a 50/50 split, you could get bias in your model

# because the observations might be affected by time e.g. first observations are higher than the last ones, etc

# so always be careful how you split your data, best way is to use "sample"



## Classification

### *KNN Classification*

```
?mtcars
```

```
attach(mtcars)
```

```
library(lattice)
```

```
# we want to create a model to classify the number of cylinders
```

```
# according to weight and mpg
```

```
with(mtcars, xyplot(wt ~ mpg, group=cyl, auto.key=T, pch=20, cex=3))
```

```
# for KNN we need to get the library class
```

```
library(class)
```

```
# lets take a look at the knn function we are using
```

```
?knn
```

```
# train: the data we use to create our model
```

```
# test: the data we use to test if the model works
```

```
# k: number of neighbors we use for the model
```

**Extra arguments:**


```
# l: a minimum amount of votes of one class
```

```
# use.all: tie handling
```

```
# prob: shows the proportion of winning class votes
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



# data frame for training

```
train <- cbind(mpg, wt)
```

# test data

```
test <- c(26, 2.2)
```

```
knn(train, test, cl = cyl, k=2, prob = T)
```

## *LDA Classification*

# we need MASS for the lda function

```
library(MASS)
```

```
?lda
```

# similar to lm and glm

```
mylda = lda(data=mtcars, cyl ~ wt + mpg)
```

```
mylda
```


# prior probabilities give the proportions of a class in the dataset

# we see the group means for each independent variable and class

# coefficients are calculated to define the areas of each class

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



```
plot(mylda)
```

```
# to use the predict function, it is useful to create a data frame with the test vectors
```

```
# test data as data frame
```

```
wt= c(2.2, 4, 1.1,5)
```

```
mpg= c(26, 20, 27, 15)
```

```
class=c(4, 6, 4, 8)
```

```
test = data.frame(wt, mpg,class)
```

```
mylda.prediction = predict(object = mylda, newdata = test[,c(1,2)])$class
```

```
# we specify that we want the class as output
```

```
mylda.prediction
```

```
table(mylda.prediction, test[,3])
```

```
# the diagonal shows the correct predictions
```

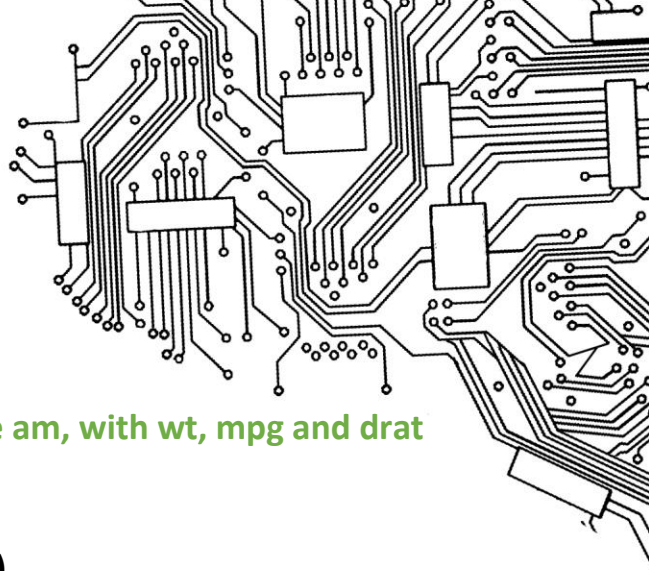
```
# as you can see in the table, there is one error in the result
```

```
# observation 2 in the test set was classified as 8
```

```
# in the data frame it is class 6
```

```
# on the plot we can see that it is quite near to class 8
```





## *Logistic Regression*

# in this case we want to model a binary outcome am, with wt, mpg and drat

```
head(mtcars)
```

```
mymtcars=data.frame(am = as.factor(mtcars$am),
```

```
    wt = mtcars$wt,
```

```
    mpg = mtcars$mpg,
```

```
    drat = mtcars$drat)
```

# since we are performing a logistic regression on a classification, we check if our outcome

# variable is in deed a factor (class)

```
class(mymtcars$am)
```

# glm with family = binomial is the classic way of logistic regression in R

```
mylog = glm(data = mymtcars, am ~ wt + mpg + drat, family = "binomial")
```

```
summary(glm(data = mymtcars, am ~ wt + mpg + drat, family = "binomial"))
```

# in this case I decide to keep all three predictors in the model

# we are going to run the model on the training data itself

```
testprediction <- predict(mylog, type="response")
```

```
testprediction
```

# prob <= 0.5 means 0 or automatic

# we can get a character vector of the 2 transmission types

```
predicted.classes = rep( "automatic",32)
```


```
predicted.classes[testprediction > .5]="manual"
```

```
predicted.classes
```



# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



```
table(predicted.classes, mymtcars$am)
```

```
# the table tells us that we had 2 misclassifications
```

```
# now we see what the model would predict for our test add-on
```


```
# for the predict function it is best to use a data frame
```

```
addon = data.frame(wt = 4.500, mpg = 30.2 , drat = 4.9)
```

```
predict(mylog, addon, type="response")
```

```
# type response for probabilities
```

```
## the model would predict that a car of 4500 lb has 0 % probability of having a  
manual transom
```



## *Exercise Iris - LDA and KNN*

# Classify the iris dataset according to Species

# using the predictors Petal.Length and Petal.Width

# perform both LDA and KNN (k of 3)

# test dataframe

```
Petal.Width = c(0.7, 2.5)
```

```
Petal.Length = c(2.4, 7)
```

```
Species = c("setosa", "virginica")
```

```
test = data.frame(Petal.Width, Petal.Length, Species)
```

```
library(lattice)
```

```
with(iris, xyplot(Petal.Length ~ Petal.Width, group=Species,  
                  auto.key=T, pch=20, cex=3))
```

```
library(MASS)
```

```
mylda = lda(data=iris, Species ~ Petal.Length+Petal.Width)
```


```
mylda
```

```
plot(mylda)
```

# to use the predict function, it is useful to create a data frame with the test vectors

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



```
# test dataframe
```

```
Petal.Width = c(0.7, 2.5)
```

```
Petal.Length = c(2.4, 7)
```

```
Species = c("setosa", "virginica")
```

```
test = data.frame(Petal.Width, Petal.Length, Species)
```

```
mylda.prediction = predict(object = mylda, newdata = test[,c(1,2)])$class
```

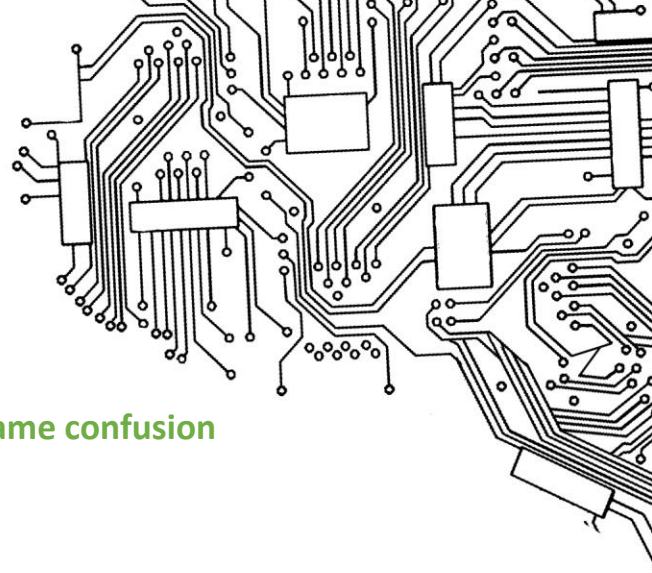
```
# we specify that we want the class as output
```

```
mylda.prediction
```

```
table(mylda.prediction, test[,3])
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



## *KNN*

# clear the environment at first to avoid object name confusion

```
attach(iris)
```

```
train <- cbind(Petal.Width, Petal.Length)
```


```
library(class)
```

```
test = matrix(c(0.7, 2.5, 2.4, 7), nrow=2)
```

```
knn(train, test, cl=Species, k=3, prob=T)
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



## Tree Based Models

# for our example we are using the tree library (one of the main packages for trees)

```
library(tree)
```

# lets work with the standard mtcars dataset

```
attach(mtcars)
```

```
?tree
```

# lets plot the data to see what we are dealing with

```
library(lattice)
```

```
with(mtcars, xyplot(mpg ~ wt, group=am, auto.key=T, pch=20, cex=3))
```

# the function is quite similar to lm

```
mytree = tree(data=mtcars, am ~ wt + mpg)
```

# as you can see the terminal nodes give us group means

# which is not useful for a Yes vs No question like transmissions


```
plot(mytree)
```

```
text(mytree)
```

```
title("Pseudo-Regression Tree MTCARS", sub="automatic vs manual  
Transmission")
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



# for this true classification tree I am factorizing am

# you can see it as qualitative now

```
mytree = tree(data=mtcars, as.factor(am) ~ wt + mpg)
```

# a summary gives a first overview on the tree

```
summary(mytree)
```

```
plot(mytree)
```

```
text(mytree)
```

```
title("Classification Tree MTCARS", sub="automatic vs manual Transmission")
```

# the terminal nodes appear to be the same, so the whole last split could have been omitted

# however it can be useful because as we can see in the previous tree, the means on the terminal nodes are very different

# lets split the dataset in half and calculate the test error rate

# I am creating a new data frame with am as factor

```
mtcars.new = data.frame(am.new = as.factor(am) ,wt, mpg)
```


```
head(mtcars.new)
```

```
class(mtcars.new$am.new)
```



# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



```
# training and test data made out of the initial mtcars.new
```

```
train = mtcars.new[1:16,]
```

```
test = mtcars.new[17:32,]
```

```
# tree made from the training set
```

```
mytree.train = tree(am.new ~ wt + mpg, data=train)
```

```
# now we run a prediction on the test data
```

```
mytree.pred = predict(mytree.train, test, type="class")
```

```
# type class for classification
```

```
# and here we compare the results predicted vs reality
```


```
table(mytree.pred, test$am.new)
```

```
treetable = table(mytree.pred, test$am.new)
```

```
(sum(diag(treetable)))/16
```

```
# in this case the tree would be correct in 87.5 % of the cases
```





## *Exercise Classification Tree*

### Example with diamonds data

- Library: ggplot2, dataset: diamonds
- create a tree to classify for color with price and the variable x
- use the first 500 rows for your tree
- plot and visualize the tree
- check the test error by splitting the set of 500 in 2 subgroups

```
library(tree)
```

```
library(ggplot2)
```

```
attach(diamonds)
```

```
mytree = tree(data=diamonds[1:500,], color ~ price + x)
```

```
summary(mytree)
```

```
plot(mytree)
```

```
text(mytree)
```

```
title("Classification Tree Color of Diamonds")
```

**# Lets check the test error rate**


```
diamonds.df = data.frame(color, price, x)
```

```
diamonds.new = diamonds.df[1:500,]
```

```
head(diamonds.new); class(diamonds.new)
```

```
train = diamonds.new[1:250,]
```

```
test = diamonds.new[251:500,]
```



```
mytree.train = tree(color ~ price + x, data=train)
mytree.pred = predict(mytree.train, test, type="class")
# type class for classification
table(mytree.pred, test$color)
(sum(diag(table(mytree.pred, test$color))))/250
```


### *Random Forests and Bagging*

- 2 methods to reduce variance in the model
- idea bagging: use bootstrapping to sample 100s of training sets and
- calculate a tree model for each of the sets - average the models
- idea randomFo: same as bagging but only a limited number of predictors is
- used to calculate a given split

```
library(randomForest)
library(ggplot2)
set.seed(123)
# mtr determines the number of predictor variables to be used
# the full number makes for a bagging approach
bagging = randomForest(formula = color ~ . , data = diamonds[1:500,], mtr=9)
plot(bagging)
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



# now it is a random forest with 3 predictors per split

```
randomFor = randomForest(formula = color ~ . , data = diamonds[1:500,],  
mtr=3)
```

# we can check the importance of the predictors

```
importance(randomFor)
```

# and we can visualize it

```
varImpPlot(randomFor)
```

# lets test the random forest with a test data frame

```
test = diamonds[501:800,]
```

```
predicted.bagging = predict(newdata=test, bagging, type = "class")
```

```
predicted.randomFor = predict(newdata=test, randomFor, type = "class")
```

```
table(predicted.bagging, test$color)
```


```
table(predicted.randomFor, test$color)
```

# we can calculate the percentage of correct predictions

```
sum(diag(table(predicted.bagging, test$color)))/300
```

```
sum(diag(table(predicted.randomFor, test$color)))/300
```

# keep in mind that color has 7 levels - 0.35 better than pure chance of 1/7



## Clustering

### *K means clustering*

# for a simple example we can use the rivers dataset

```
plot(rivers)
```

# the function is called kmeans, number of clusters = center

# nstart specifies the number of random sets to start with

```
kclust = kmeans(rivers, centers = 3, nstart = 30)
```

```
kclust
```

# we can get a visual impression of our clustered data

```
plot(rivers, col = kclust$cluster)
```

### *Hierarchical Clustering*

# simple example of Euclidean distance

# Square Root of Sum of Squares of Differences in Attributes

```
a = mtcars[1,]
```

```
b = mtcars[11,]
```


```
dist(rbind(a,b))
```

# lets get the distance matrix of the first 16 obs at first

```
dm = dist(as.matrix(mtcars[1:16,]))
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



# use the hclust function for hierarchical clustering

```
hcluster = hclust(dm) # standard method is complete linkage
```

# dendrogram

```
plot(hcluster)
```

## *Exercise K Means Clustering*

- data for the exercise: extract the first 3 numeric columns from the iris dataset
- head(iris)
- check if the extraction worked (hint: head, summary, nrow, class, ...)
- perform K means clustering on the dataset (experiment with the number of K)
- create a vector with all observation numbers of a specific cluster (hint: which, on cluster 3)
- visualize your results in a 3d plot (hint: library rgl, plot3d)

# at first lets get the dataset to work with

```
clusterdata = data.frame(iris$Sepal.Length, iris$Sepal.Width, iris$Petal.Length)
```

# checking if the data frame has the attributes we want

```
head(clusterdata); summary(clusterdata); nrow(clusterdata); class(clusterdata)
```

# I am performing 3 kmeans analyses with K of 3, 5, 8


```
clusterk3 = kmeans(clusterdata, centers = 3, nstart = 35)
```

```
clusterk5 = kmeans(clusterdata, centers = 5, nstart = 35)
```

```
clusterk8 = kmeans(clusterdata, centers = 8, nstart = 35)
```

# STATISTICAL MODELING WITH R EXAMPLES

BY R-TUTORIALS.COM



**# Lets take a look at the 3 objects**

```
clusterk3; clusterk5; clusterk8
```

**# extracting a vector with all observation IDs in a given cluster**

```
cluster3vector = which(clusterk3$cluster == 3)
```

```
cluster3vector
```

**library(rgl) # for an easy 3d scatterplot function**

**# using the plot3d function to get a 3 dimensional scatterplot**

```
plot3d(clusterdata, size = 6, col = clusterk3$cluster,  
       xlab = "", ylab = "", zlab = "", sub = "3 Clusters")
```

```
plot3d(clusterdata, size = 6, col = clusterk5$cluster,  
       xlab = "", ylab = "", zlab = "", sub = "5 Clusters")
```

```
plot3d(clusterdata, size = 6, col = clusterk8$cluster,  
       xlab = "", ylab = "", zlab = "", sub = "8 Clusters")
```

**# which plots shows be most overlapping?**